

PRIVATE / LOCAL / EVIDENCE-GROUNDED

ClaimCourt

A Private Local AI Evidence Court

Turns fuzzy human recollections and disputed commitments into cited local evidence, contradiction timelines, and approval-gated briefs.

TEAM

ZI FEI YU

**AMD AI DevMaster
Hackathon 2026**

Track 2

Development & Local Deployment
of Private AI Agents

0 CLOSED MODEL APIs

RADEON + ROCm + vLLM

A REAL WORKSPACE PROBLEM

People remember meaning, not filenames

ClaimCourt combines fuzzy private memory retrieval with an evidence-grounded decision workflow.

The screenshot displays the ClaimCourt application interface. At the top, a navigation bar includes the 'CLAIMCOURT' logo, a 'LOCAL ONLY' indicator, and system status 'RADEON GPU READY' and 'ROCm 7.2.1'. A menu bar contains 'Case Desk', 'Intent', 'Evidence', 'Courtroom', and 'Security'. The main content area is titled 'NEW LOCAL INVESTIGATION' and features a search bar with the placeholder 'What should we investigate?'. Below the search bar, there are buttons for 'Contract SLA Review', 'Find Fuzzy Memory File', and 'Scan Sensitive Secrets', along with a 'RUN PRIVATE INVESTIGATION' button. A pipeline diagram shows four stages: 'STAGE 1 UNDERSTANDING', 'STAGE 2 RETRIEVING', 'STAGE 3 COURT', and 'STAGE 4 EXPORT'. The bottom section contains three panels: 'Intent Compiled' (not compiled), 'Evidence Sources' (Local Evidence Packet), and 'Judge Verdict' (Index a private workspace to begin). A 'CURRENT CASE SUMMARY' section at the bottom states: 'The frontend is connected to the local ClaimCourt API. No evidence is fabricated before indexing.'

01

Find the right private file from an incomplete recollection.

02

Challenge disputed claims with cited supporting and opposing evidence.

03

Locate sensitive records without revealing the secret value.

FUZZY INTENT COMPILER

Do not send vague language directly to a vector store

Compile human memory into explicit constraints, relations, exclusions, and clarification policy first.

HUMAN RECOLLECTION

"Find the deck I wrote
after the customer
launch slipped - it
mentioned Q4 risk."

**LOCAL INTENT PLAN**

```
artifact: PPTX
topics: customer delay,
delivery risk
time: Q4
relation: causal_after
mode: single
clarify: false
```

**RETRIEVAL DECISION**

Final version selected
Draft grouped as
alternate
Page / slide evidence
retained
Low confidence -> clarify

320 / 320

frozen live queries passed

100%

single-target Top-1

100%

clarification accuracy

0

unsafe wrong auto-selections

EVIDENCE COURT

One local model. Three constrained roles. One citation boundary.

The workflow is auditable orchestration, not performative multi-agent autonomy.

STAGE 1

PROSECUTOR

Strongest supporting evidence



STAGE 2

DEFENSE

Counter-evidence, exceptions
and contradictions



STAGE 3

JUDGE

Schema-valid verdict using
allowed citation IDs only


HARD OUTPUT CONTRACT

```
claim / supported | contradicted | insufficient_evidence / confidence / citations / contradictions / timeline  
/ missing evidence / next action
```

PRIVATE LOCAL ARCHITECTURE

Every sensitive component stays inside the selected host

Evidence is retrieved before inference, and every downstream action is permission- and citation-gated.



LOOPBACK MODEL ENDPOINTS

127.0.0.1:8000/v1 judge + LoRA router 127.0.0.1:8001/v1 BGE embeddings

0 external model calls

PRIVACY IS A CONTROL PLANE

ClaimCourt does not rely on a privacy promise alone

The architecture blocks unsafe endpoints, raw-secret disclosure, invented citations, silent fallback and unapproved writes.

LOCAL ENDPOINT POLICY

Loopback-only by default; private hosts require explicit owner allow-listing.

SECRET REDACTION

Raw credentials are removed before model, UI, history and export paths.

CITATION ALLOW-LIST

The judge may cite only evidence IDs from the current retrieved packet.

FAIL CLOSED

Real workspaces never receive a fabricated deterministic answer when vLLM fails.

APPROVAL GATE

No local decision brief is written until the user approves the operation.

AUDIT LEDGER

Source, evidence and report identities use SHA-256 for later verification.

VERIFIED RADEON STACK

Qwen3-14B + LoRA + BGE + reranker on one W7900-class GPU

The complete local stack peaked at approximately 45.3 GB of 51.52 GB VRAM.

Qwen3-14B

BF16 judge and controlled roles

gfx1100

AMD Radeon PRO W7900 class

45.3 GB

observed peak stack VRAM

8,192

judge context tokens

WHY THIS FITS

- One shared judge service for all three roles
- Small dedicated BGE pooling runner
- Shortlist-only cross-encoder reranking
- Evidence packet instead of full-workspace context

VERIFIED LOCAL ENDPOINTS

8000 / Qwen3-14B + ClaimCourt LoRA
8001 / BGE-small pooling runner
8503 / local ClaimCourt API
8502 / private browser interface

LIVE, NO-FALLBACK RESULTS

The quality gate requires the complete local stack

A component that fails or silently falls back causes the Radeon holdout to fail.

0.27 s

first-token latency

14.3

generation tokens / second

126.44 s

complete three-role latency

1,778

completion tokens

320 / 320

frozen fuzzy-intent holdout

100%

Top-1 / coverage / no-match / clarification

0

unsafe wrong auto-selections Router + BGE + reranker

LIVE

FOUR-MINUTE JUDGE JOURNEY

Show the dispute, the evidence, the permission gate and the GPU proof

Use only the synthetic championship workspace during public recording.

0:00	DISPUTE	Sales promise vs signed contract
0:20	RADEON	rocminfo / vLLM / local endpoints
0:45	INTENT	Compile the vague claim
1:10	EVIDENCE	Retrieve signed, draft and contradictory records
1:45	COURT	Prosecute, defend and judge
3:05	APPROVE	Permission-gated local brief
3:30	PROOF	Latency, throughput, VRAM and 320/320

RECORDING RULES

Use synthetic files only
Warm the model once
Show local vLLM, never
fallback
Hide keys, tokens and shell
history
Do not expose a real
workspace publicly
End on the measured evidence

CLAIMCOURT

Private evidence before confident answers.

ClaimCourt understands fuzzy human memory, retrieves the right private records, exposes contradictions, and asks permission before it acts.

WHY IT MATTERS

LOCAL
AUDITABLE
CITED
PRIVACY-GATED
RADEON-
OPTIMIZED